

# Clinical Skin Lesion Diagnosis

Madamanchi Sai Venkat, Vemula Giri Manohar, Vemavarapu Tejashwini

## 1. Abstract

Deep Learning and Convolutional Neural Networks are used to diagnose skin lesions in this paper. In modern healthcare, skin cancer detection through advanced image analysis is a major issue. This research classifies dermatoscopic images of skin lesions into seven categories, each indicating a different skin cancer. The study addresses data imbalance using the "HAM10000" dataset of diverse skin lesion images. This imbalance, common in healthcare datasets, has large class representation gaps. This study trains deep learning models with an adjusted emphasis to reduce class overrepresentation. The research methodology uses a custom, simple but effective CNN and a more complex VGG-like model. To balance simplicity and efficiency, the custom CNN has convolutional, max-pooling, and fully connected layers. CNNs are used for medical image analysis, but this project compares architectural models. The findings could help dermatologists screen and diagnose skin cancers early and improve treatment.

## 2. Introduction

### 2.1. Overview of the Project

Skin cancer is one of the most common cancers worldwide, making diagnosis a major public health issue. Utilizing dermatologic oncology advances like digital imaging and machine learning is the goal. Revolutionary, non-invasive methods for early skin lesion detection and classification have emerged from these technologies. We want to use these advances to improve skin cancer diagnosis accuracy and efficiency.

### 2.2. Motivation

The need for accurate and timely skin cancer diagnoses drove this project. A precise diagnosis is crucial to effective treatment and better patient outcomes. Traditional methods are laborious and require expertise. A Deep Learning system using Convolutional Neural Networks (CNNs) to assist dermatologists addresses these challenges. This system could drastically improve diagnostic efficiency and reduce subjective expertise.

### 2.3. Our Approach

Our approach is situated at the intersection of medical imaging and artificial intelligence. We aim to harness the capabilities of CNNs for diagnosing skin lesions. By utilizing a comprehensive dataset of dermatoscopic images, the project explores the potential of Deep Learning algorithms in classifying various types of skin lesions. This approach not only applies the latest technological advancements in AI to healthcare but also includes a comparative study of different CNN architectures to determine the most effective method in medical image analysis.

### 2.4. Dataset for Experiments and Evaluation

For our experiments and evaluation, we are utilizing the *HAM10000* dataset. This dataset is a large collection of dermatoscopic images, providing detailed visual information crucial for identifying different types of skin cancers. By employing this dataset, we aim to train and test our Deep Learning models, ensuring a robust evaluation of their capabilities in classifying skin lesions accurately.




FIGURE 1. Some Random images from HAM10000

### 3. Background

The HAM10000 dataset has emerged as a critical asset in the field of machine learning for skin lesion analysis, providing a vast array of dermatoscopic images from diverse populations. Esteemed in academic circles, its over 10,000 images with specific identifiers have been instrumental in developing and refining diagnostic models in dermatology. Utilized extensively in prior studies, this dataset facilitates the training of algorithms to discern various skin lesion types, with its broad age range of patients (from infants to the elderly, averaging 51.86 years) adding to its richness, despite some missing age data.

Past research has underscored the importance of the dataset's varied diagnostic categories, sourced from four diagnostic methods including histopathology, covering seven different skin conditions. This aspect, combined with the inclusion of diverse genders and the localization of lesions across 15 body sites, offers a realistic framework for model training and assessment. The dataset's detailed demographic information, highlighting a slight male predominance and the distribution of lesions across different body parts, has been pivotal in developing models that are adaptable to various lesion appearances and skin textures, enhancing the efficacy of machine learning in dermatological imaging and diagnostics.

This variety ensures that the model developed will be exposed to a wide range of skin lesions, enhancing its diagnostic capabilities.




FIGURE 2. Lesion localization over 15 anatomical regions

### 4. Approach


#### 4.1. Data Preparation and Preprocessing

- Dataset Acquisition:** Utilization of the HAM10000 dataset, comprising 10,015 dermatoscopic images across seven skin cancer categories. In analyzing the dataset, each lesion is uniquely identified, with a total of 7,470 distinct lesions, indicating that some images may be of the same lesion.
- Data Preprocessing:**
  - Image Processing:* Images are resized, normalized, and potentially augmented for neural network training.
  - Handling Class Imbalance:* Application of oversampling or similar techniques to address class imbalances within the dataset.
  - Analysing the distribution:* The localization and demographic information will be used to explore the model's performance across different patient groups and lesion sites, which could provide valuable insights into the model's applicability in a clinical setting

4.2. Model Development

a. Model Architecture:

- Custom CNN Model: Sequential architecture with convolutional layers, max pooling, flattening, and dense layers using ReLU and softmax activations.
- VGG-Inspired Model: A deeper CNN model inspired by the VGG architecture, known for its efficacy in image classification.



- **Compilation:** Models compiled with sparse categorical cross-entropy loss and Adam optimizer.
- **Callbacks:** Use of ModelCheckpoint to monitor validation accuracy and save the best performing model.

4.3. Model Training

- a. **Training Process:** Training of the models for a defined number of epochs with a set batch size, including a validation subset to monitor performance and mitigate overfitting.

4.4. Model Evaluation and Results

- a. **Performance Visualization:** Plotting of loss and accuracy against epochs for both training and validation phases.
- b. **Model Testing:** Evaluation of the models on a separate test set to determine classification accuracy.
- c. **Results Analysis:** Analysis of training and testing outcomes to assess model effectiveness.

4.5. Additional Steps

- *Hyperparameter Tuning:* Optimization of model performance through adjustments in learning rate, layer count, or neurons per layer.
- *Regularization Techniques:* Implementation of dropout, L2 regularization, or other methods to prevent overfitting.
- *Data Augmentation:* Potential enhancement of the model's generalization capability through data augmentation.

5. Results

5.1. Dataset Characteristics

The HAM10000 dataset showcased a significant imbalance in class distribution among the different skin lesions, with Melanocytic nevi (nv) being the most prevalent and Dermatofibroma (df) the least. This imbalance poses a challenge for model training as it can lead to a bias towards more frequently represented classes. Gender distribution within the dataset was relatively balanced with a slight male predominance, which is consistent with the general demographics of skin cancer incidence. The disease distribution across body locations indicated that certain skin lesions were more common in specific body areas, offering potential insights into lesion localization that could be valuable for improving the specificity of the CNN models.




FIGURE 4. Imbalanced Data




FIGURE 5. Balanced Data

## 5.2. Data Visualization and Preprocessing

The paper commenced with an analysis of the class distribution within the HAM10000 dataset, highlighting the challenges of data imbalance common in medical datasets. To mitigate this, oversampling techniques were employed, resulting in a more balanced dataset for model training. Furthermore, patient demographics were explored, with a pie chart visualization illustrating the gender distribution, ensuring the model's broad applicability.

## 5.3. Training and Validation Performance

The training outcomes for the two CNN models were as follows:

- The custom-built CNN model attained a training accuracy of 81.51% and a validation accuracy of 76.23%. However, an upward trend in validation loss during later epochs was observed, indicating potential overfitting.
- The VGG-like model exhibited a training accuracy of 89.67% and a validation accuracy of 75.11%. Similar to the first model, its validation loss trends suggested overfitting.
- Evaluation on the test set yielded accuracies: 61.41% for the custom CNN model and 56.27% for the VGG-inspired model. This points to moderate performance on unseen data and highlights the discrepancy between training/validation and test accuracy.

Graphs plotting the loss and accuracy against epochs provided visual insights into the models' performance over time.



FIGURE 6. Simple CNN Model Accuracy and Loss




FIGURE 7. VGG Model Accuracy and Loss

## 5.4. Dataset Evaluation and Generalization

Once the Convolutional Neural Network (CNN) model was trained, it was employed to label images from the dataset, which served as a crucial step in the evaluation process. The model utilized its learned weights and features to make predictions on the test set, consisting of images that were not exposed to the model during training. Each dermatoscopic image was processed through the model's multiple layers, resulting in a classification output that assigned a label corresponding to one of the seven skin lesion categories present in the HAM10000 dataset. This suggests a need for further tuning, possibly through enhanced regularization methods or data augmentation, to improve the models' generalization capabilities.

## 5.5. Loss Trends and Overfitting Concerns

The loss trends highlighted a common challenge in machine learning: as the model became more adept at classifying the training data (decreasing training loss), it did not show a proportional decrease in validation loss. This discrepancy suggested that while the model was memorizing the training data, it was not learning the underlying patterns necessary for robust generalization. These results underscore the importance of implementing strategies to mitigate overfitting, such as data augmentation, dropout layers, or more complex regularization techniques.

In conclusion, while the models demonstrated effective learning from the training data, their performance on the test set indicates a need for ongoing refinement. The disparity between training/validation and test accuracy underscores the importance of developing models that generalize well to new data, a crucial aspect in the field of medical diagnostics and AI-driven healthcare.




FIGURE 8. Prediction of images using trained model

## 6. Discussion

### 6.1. Model Performance Interpretation and Challenges

The custom CNN and VGG-inspired models showed varying success in classifying seven skin cancer types. A key challenge identified was the impact of data imbalance on model performance, with common lesion types being overrepresented. This could reduce sensitivity for rarer lesions like melanoma. To counteract this, a more balanced dataset and techniques such as oversampling or advanced data augmentation are suggested. Additionally, the choice of CNN architecture is crucial, as deeper networks might overfit on limited datasets.

### 6.2. Clinical Implications and Future Directions

Enhancements in CNN-based dermatological diagnostic systems are proposed, including using larger, more diverse datasets and exploring advanced architectures like ResNet or Inception models. These improvements could enhance accuracy and generalizability. However, these AI models should complement, not replace, medical judgment. The integration of AI in clinical practice requires careful consideration of technical, ethical, and workflow aspects to ensure reliable support for dermatologists. By refining these

models and carefully integrating them into clinical practice, we can realize AI's potential to improve skin cancer detection and diagnosis, improving patient outcomes.

## 7. Conclusion

In conclusion, the deployment of Convolutional Neural Networks (CNNs) for the classification of skin lesions from the HAM10000 dataset has showcased the potential of deep learning in enhancing diagnostic processes in dermatology. The models trained in this study have capably handled a diverse and imbalanced dataset, learning to discern intricate patterns within the dermatoscopic images. Despite the challenges posed by class imbalances and the high variability of skin lesions, the CNNs achieved a commendable level of accuracy on the test set. This indicates a foundational success in the models' ability to generalize from their training and suggests a significant step forward in the application of artificial intelligence in medical image analysis.

However, the project also highlighted critical areas for improvement, particularly in model generalization and overfitting. The observed discrepancies between training accuracy and validation performance signal the need for enhanced regularization techniques and model tuning to ensure consistency across unseen datasets. As the field progresses, further refining these models will be essential to support their clinical adoption. The ultimate goal is to develop an AI-assisted diagnostic tool that provides reliable and robust support to dermatologists, contributing to early and accurate skin cancer diagnosis and thereby improving patient care outcomes.

## A. References

1. P. Tschandl, C. Rosendahl, and H. Kittler, "The HAM10000 dataset, a large collection of multi-source dermatoscopic images of common pigmented skin lesions," *Scientific Data*, vol. 5, no. 1, Aug. 2018. [Online]. Available: <http://dx.doi.org/10.1038/sdata.2018.161>
2. N. Codella, V. Rotemberg, P. Tschandl, M.E. Celebi, S. Dusza, D. Gutman, B. Helba, A. Kalloo, K. Liopyris, M. Marchetti, et al., "Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration (ISIC)," arXiv preprint arXiv:1902.03368, 2019.
3. P. Tschandl, "HAM10000 dataset [Data set]," Harvard Dataverse, V1, 2018. [Online]. Available: <https://doi.org/10.7910/DVN/DBW86T>

## B. Appendix

### B.1. Model Equations and theories

The primary mathematical operations in a CNN are convolutions, non-linear activation functions, pooling, and fully connected layers. The convolution operation is defined as:

$$(F * I)(x, y) = \sum_m \sum_n F(m, n) \cdot I(x - m, y - n)$$

where  $F$  is the filter,  $I$  is the input image, and  $x, y$  are the coordinates in the image and filter respectively.

The Rectified Linear Unit (ReLU) function introduces non-linearity into the model and is defined as:

$$ReLU(x) = \max(0, x)$$

Max pooling is a common pooling operation which reduces the spatial dimensions and is defined as:

$$MaxPooling(R) = \max_{x \in R} x$$

where  $R$  is the pooling region.

The softmax function, used for multi-class classification, provides probabilities for each class and is defined as:

$$Softmax(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

for class scores  $z_i$  from the output of the network.

The cross-entropy loss function for multi-class classification is defined as:

$$CrossEntropy(y, p) = - \sum_i y_i \log(p_i)$$

where  $y$  are the true labels and  $p$  are the predicted probabilities.

The mathematical representation of a dense layer is:

$$y = f(W \cdot x + b)$$

where  $x$  is the input vector,  $W$  is the weight matrix,  $b$  is the bias vector, and  $f$  is the activation function.

The Adam optimizer combines the best properties of the AdaGrad and RMSProp algorithms. Adam optimizer updates parameters based on the adaptive estimation of first and second-order moments.

Overfitting is a condition when a model learns the training data too well, including noise and outliers. Regularization techniques such as dropout or L2 regularization can be employed to prevent this.

### B.2. Experiments, Evaluation, and Visualization of the Dataset




FIGURE 9. Class Distribution of HAM10000




FIGURE 10. Gender Distribution of HAM10000




FIGURE 11. Distribution based on Patients age